

# **Delta lake format introduction**

# Disclaimer (I)

Most of the code is based on jupyter notebooks using the spylon kernel (scala), but there may be some .dbc files

# Disclaimer (II)

You can enable delta on your local spark by download the jar and enabling the extensions.

```
bin/spark-shell --packages io.delta:delta-spark_2.12:3.0.0  
--conf "spark.sql.extensions=io.delta.sql.DeltaSparkSessionExtension"  
--conf "spark.sql.catalog.spark_catalog=org.apache.spark.sql.delta.catalog.DeltaCatalog"
```

# Disclaimer (III)

<https://docs.delta.io/latest/delta-standalone.html>

# Disclaimer (IV)

There is a docker-compose file in the project granting you access to a jupyter notebook with delta

# What is Delta Lake?

It was a project launched by [Databricks in 2017](#) for merging the data warehouse and the datalake and it become managed by the Linux Foundation in 2019.

# Advantages?

- ACID transactions on file formats
- Efficient data layout
- Rename, reorder columns in a table
- Metadata scale
- Time travel to older versions of the table
- Enable change data feed aka change data capture

<https://www.vldb.org/pvldb/vol13/p3411-armbrust.pdf>

## **How does it work?**

It is based on parquet files containing the actual data + json files describing a protocol.



**Demo time**

# From Parquet to Delta

```
import io.delta.tables._  
  
DeltaTable.convertToDelta(spark, "parquet.`sample_data/user_data`")
```



tree files delta

# Delta from a dataframe

```
val sampleData = Seq(  
  (1, "Pedro"),  
  (2, "Sergi"),  
  (3, "Aleix"),  
  (4, "David")  
)  
  
import spark.implicits._  
val df = sampleData.toDF()  
df.coalesce(1).write.mode("overwrite").format("delta").saveAsTable("data_engineers")
```

# Table directory

```
spark-warehouse/data_engineers/
|->delta_log/
|          |->00000000000000000000000000000000.json
|->part-00000-1642cfd8-1dfa-497a-8c3f-2e000c2b1d79-c000.snappy.parquet
```

# Delta Log (I)

```
{
  "commitInfo": {
    "timestamp": 1698680849935,
    "operation": "CREATE OR REPLACE TABLE AS SELECT",...
  }
}
{
  "protocol": {
    "minReaderVersion": 1,
    "minWriterVersion": 2
  }
}
{
  "metaData": {
    "id": "ab71b946-f933-4dff-a11f-e7c6ecf70264",
    "format": {
      "provider": "parquet",
      "options": {}
    },
    "partitionColumns": [],
    "configuration": {},
    "createdTime": 1698680849740
  }
}
```

# Delta Log (II)

```
{
  "add": {
    "path": "part-00000-1642cfd8-1dfa-497a-8c3f-2e000c2b1d79-c000.snappy.parquet",
    "partitionValues": {},
    "size": 714,
    "modificationTime": 1698680849845,
    "dataChange": true,
    "stats": "{\"numRecords\":4,\"minValues\":{\"_1\":1,\"_2\":\"Aleix\"},\"maxValues\":{\"_1\":4,\"_2\":\"Sergi\"},\"nullCount\":{\"_1\":0,\"_2\":0}}\"
  }
}
```

# Merging data (I)

```
import io.delta.tables._
```

```
val newDataDf = Seq(  
  (4, "Adrian"),  
  (5, "David"),  
  (6, "Simon")  
) .toDF()
```

```
DeltaTable
```

```
.forName("data_engineers")  
.as("oldData").merge(newDataDf.as("newData"), "oldData._1 = newData._1")  
  .whenMatched(some condition)  
  .updateAll  
  .whenNotMatched  
  .insertAll  
  .execute()
```

## **How does it get atomicity for a merge?**

1. Find files in the table that match join condition
2. Read those files and write them again with updates and / or inserted rows
3. Add a new entry in the protocol removing the touched files and adding the new ones



# Merging data (II)

## Merge data (III)

insertOnly -> leftAntiJoin (left join where right is null)

updateOnly -> rightJoin

insert, update and delete -> fullOuterJoin

## Merging data (III)

```
spark-warehouse/data_engineers/
|->delta_log/
|           |->00000000000000000000000000000000.json
|           |->00000000000000000000000000000000.json.crc
|           |->00000000000000000000000000000001.json
|           |->00000000000000000000000000000001.json.crc
|->part-00000-1642cfd8-1dfa-497a-8c3f-2e000c2b1d79-c000.snappy.parquet
|->part-00000-1642cfd8-1dfa-497a-8c3f-2e000c2b1d79-c000.snappy.parquet.crc
|->part-00000-17fb70c7-4412-4a3e-8e03-d12d0e41b449-c000.snappy.parquet
|->part-00000-17fb70c7-4412-4a3e-8e03-d12d0e41b449-c000.snappy.parquet.crc
```

# Merging data (IV)

```
{
  "commitInfo": {
    "timestamp": 1698682774893,
    "operation": "MERGE",
  },
  "remove": {
    "path": "part-00000-17fb70c7-4412-4a3e-8e03-d12d0e41b449-c000.snappy.parquet",
    "deletionTimestamp": 1698682774892,
    "size": 726
  },
  "add": {
    "path": "part-00000-c906bcfb-2fe9-49bc-b07b-c65cc9b5fdcc-c000.snappy.parquet",
    "modificationTime": 1698682774892,
    "dataChange": true,
    "stats": "{\n  \"numRecords\":5,\n  \"minValues\":{\n    \"_1\":1,\n    \"_2\":\n      \"Adrian\"\n  },\n  \"maxValues\":{\n    \"_1\":5,\n    \"_2\":\n      \"Simon\"\n  },\n  \"nullCount\":{\n    \"_1\":0,\n    \"_2\":0\n  }\n}"
  }
}
```

# Delta Log (III)

# Checkpoints and last\_checkpoint

```
5.to(14).foreach(i => {
    val newIntern = Seq((i, s"Intern ${i}")).toDF("id", "engineer")
    newIntern.write.mode("append").format("delta").saveAsTable("data_engineers_checkpoint")
})
```

```
spark-warehouse/data_engineers_cdc/  
|->_delta_log/  
|      |->_last_checkpoint  
|      |->000000000000000000000000.json  
|      |->...  
|      |->000000000000000000000010.checkpoint.parquet
```

# Delta Log (IV)

```
{  
  "version": 10,  
  "size": 13,  
  "sizeInBytes": 15001,  
  "numOfAddFiles": 11,  
  "checkpointSchema": ...  
}
```

txn	add	remove	metaData	protocol
null	{part-00000-d08223e8-87b6-487c-bcc5-837c64e24d2e-c000.snappy.parquet, {}, 733, 1701035863803, ... }	null	null	null

# Time travel

```
[27]: spark.read.format("delta").option("versionAsOf", 0).load("./spark-warehouse/data_engineers").show(false)
```

```
23/10/30 16:28:16 INFO Executor: Running task 0.0 in stage 117.0 (TID 825)
23/10/30 16:28:16 INFO FileScanRDD: Reading File path: file:/home/jovyan/spark-warehouse/data_engineers/part-0
23/10/30 16:28:16 INFO Executor: Finished task 0.0 in stage 117.0 (TID 825). 1698 bytes result sent to driver
23/10/30 16:28:16 INFO TaskSetManager: Finished task 0.0 in stage 117.0 (TID 825) in 9 ms on 0541434f17db (exe
23/10/30 16:28:16 INFO TaskSchedulerImpl: Removed TaskSet 117.0, whose tasks have all completed, from pool
23/10/30 16:28:16 INFO DAGScheduler: ResultStage 117 (show at <console>:43) finished in 0.011 s
23/10/30 16:28:16 INFO DAGScheduler: Job 71 is finished. Cancelling potential speculative or zombie tasks for
23/10/30 16:28:16 INFO TaskSchedulerImpl: Killing all running tasks in stage 117: Stage finished
23/10/30 16:28:16 INFO DAGScheduler: Job 71 finished: show at <console>:43, took 0.012766 s
```

```

+-----+
|_1|_2|
+-----+
|1|Aleix|
|2|David|
|3|Sergi|
|4|Pedro|
+-----+

```

# Delta Restore (I)

```
val deltaTable = DeltaTable  
  .forName("data_engineers_restore")  
  
deltaTable.restoreToVersion(0)
```



## Delta Restore (II)

```
{
  "commitInfo": {
    "timestamp": 1701005204513,
    "operation": "RESTORE",
  }
  "add": {
    "path": "part-00000-4a1eabf4-8001-4114-a8a1-3658ef390cbe-c000.snappy.parquet",
    "stats": "{...}"
  }
  "remove": {
    "path": "part-00000-89eda872-b3ea-4583-8cfa-80070de0f091-c000.snappy.parquet",
    "deletionTimestamp": 1701005204557,
    "size": 740
  }
}
```

## Limits

`delta.logRetentionDuration` = "Each time a a checkpoint is written, <Databricks> automatically cleans up log entries older than the retention interval. Default 30 days.

`delta.deletedFileRetentionDuration` = "controls how long ago a file must have been deleted before being a candidate for VACUUM. The default is interval 7 days.

# Manual Maintenance

Two main operations:

- Compaction: Reorder files for improving reading speeds
- Vacuum : Get rid of unused versions

# Compaction (I)

```
OPTIMIZE my_table
```

```
spark.read.parquet("small_files").coalesce(1).write
```

# Compaction (II) Z-order

```
OPTIMIZE my_table zorder by id
```

<https://github.com/delta-io/delta/blob/13f7fbce7b89cec387df9fbaba0389fe892322b8/spark/src/main/scala/org/apache/spark/sql/delta/expressions/InterleaveBits.scala#L81>

# Conflicts

INSERT (1)		UPDATE, DELETE, MERGE INTO	OPTIMIZE
INSERT	Cannot conflict		
UPDATE, DELETE, MERGE INTO	Cannot conflict in WriteSerializable. Can conflict in Serializable; see avoid conflicts with partitions.	Can conflict in Serializable and WriteSerializable; see avoid conflicts with partitions.	
OPTIMIZE	Cannot conflict	Can conflict	Can conflict

# Vacuum

```
VACUUM my_table DRY RUN      -- do dry run to get the list of files to be deleted  
VACUUM my_table
```

## “ **Warning**

VACUUM WILL REMOVE ALL FILES THAT ARE NOT LISTED IN THE DELTA TABLE UNLESS THEY START BY \_, if you have a checkpoint in the same path call it \_checkpoint

”

## Parallel deletes

```
spark.databricks.delta.vacuum.parallelDelete.enabled -> true
```

# Column Mapping (I)

```
spark.sql("""
  ALTER TABLE data_engineers_mapping SET TBLPROPERTIES (
    'delta.minReaderVersion' = '2',
    'delta.minWriterVersion' = '5',
    'delta.columnMapping.mode' = 'name'
  )
""")
```



# Column Mapping (II)

```
spark.sql("ALTER TABLE data_engineers_mapping RENAME COLUMN engineer TO employee")
```

```
"commitInfo": {
  "timestamp": 1701011573250,
  "operation": "RENAME COLUMN",
  "schemaString": {
    "type": "struct",
    "fields": [
      {
        "name": "id",
        "type": "integer",
        "nullable": true,
        "metadata": {
          "delta.columnMapping.id": 1,
          "delta.columnMapping.physicalName": "id"
        }
      },
      {
        "name": "employee",
        "type": "string",
        "nullable": true,
        "metadata": {
          "delta.columnMapping.id": 2,
          "delta.columnMapping.physicalName": "engineer"
        }
      }
    ]
  }
}
```

# Column Mapping (III)

```
spark.sql("ALTER TABLE data_engineers_mapping DROP COLUMN employee")
```

```
{
  "type": "struct",
  "fields": [
    {
      "name": "id",
      "type": "integer",
      "nullable": true,
      "metadata": {
        "delta.columnMapping.id": 1,
        "delta.columnMapping.physicalName": "id"
      }
    }
  ]
}
```

## Change data feed (I)

```
ALTER TABLE myDeltaTable SET TBLPROPERTIES (delta.enableChangeDataFeed = true)
```

“ WARNING: only for future changes after enabled! ”

```
org.apache.spark.sql.delta.DeltaAnalysisException:  
Error getting change data for range [0 , 2] as change data was not  
recorded for version [0].  
If you've enabled change data feed on this table,  
use `DESCRIBE HISTORY` to see when it was first enabled.
```

## Change data feed (II)

```
spark.read.format("delta")  
  .option("readChangeFeed", "true")  
  .option("startingVersion", 0)  
  .table("data_engineers_cdc").show(false)
```

_1	_2	_change_type	_commit_version
4	David	update_preimage	1
4	Adrian	update_postimage	1
5	David	insert	1

## Change data feed (III)

```
spark-warehouse/data_engineers_cdc/  
|->_change_data/  
|      |->cdc-00000-540b06a2-4586-48e4-af82-7e48ffdd89c4.c000.snappy.parquet  
|->_delta_log/  
|      |->000000000000000000000000.json  
|      |->000000000000000000000001.json  
|->part-00000-96291cc4-788d-46aa-8aa1-272cbf1e04dd.c000.snappy.parquet  
|->part-00000-b6e25d04-253f-49af-bd09-ab65916c8675-c000.snappy.parquet
```

# The future

- Deletion Vectors: mark in metadata rows as remove and rewrite the files during an optimize
- Row - level concurrency: allow incompatible operations to run together such as merge and optimizes

## **Going back to parquet**

If you want to go back from Delta to Parquet you need to:

1. Get rid of old versions of the data.
2. Clean the delta\_log folder.

## Going back to parquet (II)

```
VACUUM <table> RETAIN 0 HOURS
```

```
spark.databricks.delta.retentionDurationCheck.enabled false
```



# Extra resources

<https://www.databricks.com/blog/2019/08/21/diving-into-delta-lake-unpacking-the-transaction-log.html>

<https://www.databricks.com/blog/2020/09/29/diving-into-delta-lake-dml-internals-update-delete-merge.html>

<https://docs.delta.io/0.7.0/delta-batch.html#-data-retention>

<https://canadiandataguy.medium.com/delta-vs-parquet-a-deep-dive-into-big-data-storage-solutions-34ffab2fea52>

<https://github.com/delta-io/delta/blob/13f7fbce7b89cec387df9fbaba0389fe892322b8/spark/src/main/scala/org/apache/spark/sql/delta/commands/VacuumCommand.scala#L323>